July 18th: Virtual viewing party

4-4:30: Lightning talks for virtual posters 1

4:30-5: Virtual poster session 1

- **Daniel Hardt**. Sparks of Pure Competence in LLMs: the Case of Syntactic Center Embedding in English
- **Brandon Prickett**. Explaining differences between phonotactic learning biases in the lab and typological trends using Probabilistic Feature Attention
- **Kanishka Jain, Ashwini Vaidya**. Investigating the Probability of External Causation in Hindi Light Verb Constructions
- **Anil Korde, Philip Resnik**. On the Dangers of Naïve Replication: The Case of Implicature
- **Nathaniel Imel, Christopher Haberland, Shane Steinert-Threlkeld**. The Unnatural Language ToolKit (ULTK)
- **Giorgio Magri.** A principled derivation of OT and HG within constraint-based phonology.

5-5:30: Lightning talks for virtual posters 2

5:30-6: Virtual poster session 2

- **Lauren Levine**. A Cross-Genre Analysis of Discourse Relation Signaling in the GUM Corpus
- **Gregory M Kobele, Lei Liu**. Formalizing Feature Inheritance
- **Kayla Shames**. Designing a Digital Keyboard for Itunyoso Triqui
- **Nitin Venkateswaran, Rachel Meyer, Ratree Wayland.** Measuring the Impact of Segmental Deviation on Perceptions of Accentedness using Gradient Phonological Class Features
- **Richard Futrell**. Language Learning as Codebreaking: The Key Roles of Redundancy and Locality
- **Coleman Haley, Sharon Goldwater, Edoardo Ponti**. Visual groundedness as an organizing principle for word class: Evidence from Japanese
- **Cesare Spinoso-Di Piano, David Eric Austin, Pablo Piantanida, Jackie CK Cheung**. (RSA)^2: A Rhetorical-Strategy-Aware Rational Speech Act Framework for Figurative Language Understanding

July 19th:

9-9:30am: Opening remarks

9:30-10:30: Keynote 1: Malihe Alikhani

Theory of Mind in Generative Models: From Uncertainty to Shared Meaning

We will explore how generative models can effectively facilitate communicative grounding by incorporating theory of mind alongside uncertainty and human feedback. We begin by examining how models signal and quantify predictive uncertainty, highlighting computational parallels to epistemic stance. Next, we discuss belief modeling, presenting evidence that language models can infer degrees of interlocutor uncertainty, a crucial component in managing reference and intent. We address how a failure to accurately track beliefs may lead to sycophancy, or over-alignment with user views. We then explore the positive role of friction introduced through structured discourse or interactional pauses, which slows down interactions to promote clarity and facilitate grounding. Finally, we extend these concepts to multimodal and socially situated contexts, drawing on research in sign language modeling and human-in-the-loop training to illustrate how shared meaning can be constructed across diverse modalities and populations. This line of research demonstrates how generative models embody core mechanisms of pragmatic reasoning, offering linguists and cognitive scientists both methodological challenges and opportunities to question how computational systems reflect and shape our understanding of meaning and interaction.

10:30-11: Coffee break

11-12: Pragmatics Session 1 (Chair: Andrew Kehler)

- **Ryan Ka Yau Lai, Yan Lashchev**. Dimensions of (dis)preference in designing polar answers in American English: A latent class analysis
- **Daniil Ignatev, Denis Paperno, Massimo Poesio**. Annotator disagreement in RST annotation schemes
- **Muxuan He, Elsi Kaiser, Khalil Iskarous**. Modeling Pragmatic Accommodation: A [Rational, Irrational] Speaker in a [Wonky, Normal] World

12-1: Keynote 2: Junyi Jessy Li

Discourse Models with Questions Under Discussion

Strong generation capabilities of modern LLMs have enabled new models and investigations into the structure, processing, and understanding of discourse. In this talk, I present our efforts on developing discourse models under the Questions Under Discussion (QUD) framework, using question generation to analyze the progression of discourse. With design and training, LLMs can resurface curiosity-driven potential questions from readers, ground their elicitation in text, and connect reader expectations with writer intention. Next, I demonstrate how such generative discourse models can be used to quantify model behavior that seems intuitive and yet elusive. With QUD models, I present two measures: (a) measure for structural similarities in LLM-generated texts, which existing lexical and syntactic based similarity metrics fail to capture; (b) a measure of information salience in LLMs, using summarization as a behavioral probe.

1-2:15: Lunch (Pop Up Mentoring for students)

2:30-3:30: Talk session 1

- **Natalia Tyulina**. Empirical Analysis of Russian Aspectual Prefixes: A Computational Approach to Productivity & Semantic Relatedness
- **Hayley Ross, Kathryn Davidson, Najoung Kim**. Is analogy enough to draw novel adjective-noun inferences?
- **Gustavo Cilleruelo Calderón, Emily Allaway, Barry Haddow, Alexandra Birch**. MGen: Millions of Naturally Occurring Generics in Context

3:30-4: Coffee break

4-5:30: Talk session 2 (Chair: Canaan Breiss)

- **Edward Flemming, Giorgio Magri**. Strict domination in probabilistic phonology
- **Hyunjung Joo, Adam Jardine**. Intonation as a quantifier-free logical interpretation of metrical and prosodic structure
- **Jonathan Charles Paramore**. Learning Covert URs via Disparity Minimization
- **Caleb Belth**. A Theory of When and How Learners Construct Tiers: Implications for Opaque and Transparent Vowels

5:30-6:30: Poster session 1 (and reception)
Sponsored by Google

- **Nina Haket, Ryan Daniels.** BERT's Conceptual Cartography: Mapping the Landscapes of Meaning
- **Michael Kamerath, Aniello De Santo.** Do LLMs Disambiguate Italian Relative Clause Attachment?
- **Josephine Kaminaga, Jennie Wu, Daniel Yeung, Simon Todd**. Aligning Embedding Spaces Across Languages to Identify Word Level Equivalents in the Context of Concreteness and Emotion
- **Jonathan B. Sakunkoo, Annabella Sakunkoo**. Mind the Gap: Computational Quality Assurance of Crowd-Sourced Linguistic Knowledge on Latin and Italian Morphological Gaps
- **Iona Carslaw, Sivan Milton, Nicolas Navarre, Ciyang Qing, Wataru Uegaki**. Automatic Extraction of Clausal Embedding Based on Large-Scale English Text Data
- **Jane Li, Alan Zhou**. CNNs that robustly compute vowel harmony do not explicitly represent phonological tiers
- **Jwalanthi Ranganathan, Rohan Jha, Kanishka Misra, Kyle Mahowald**. semantic-features: A User-Friendly Tool for Studying Contextual Word Embeddings in Interpretable Semantic Spaces
- **Mary Kennedy**. Evidence of Hierarchically-Complex Syntactic Structure Within BERT's Word Representations

July 20th:

9-10: Keynote 3: Hannah Rohde

Why am I Saying This? Sensible and Informative Contributions to Discourse

For a conversation to be coherent, we expect more than just a sequence of arbitrary sentences -- we expect the sentences to connect in predictable and sensible ways. At the same time, if everything a speaker said were entirely predictable, there'd be no new information conveyed. These two expectations -- one for plausibility, one for informativity -- are two key components that underlie cooperative communication. This talk presents a series of case studies on what speakers choose to talk about and what forms they use, with

subsequent tests of how such choices in turn influence comprehenders' expectations, notably when they're estimating what words are coming next and when they're trying to recover a speaker's intended meaning.

10-10:30: Coffee break

10:30-12: Talk session 3

- **Kenneth Hanson**.Adjunction in (T)SL Syntax
- **Lei Liu**. Left-corner Minimalist parsing of mixed word order preferences
- **Logan Swanson, Kenneth Hanson, Thomas Graf**. Are syntactic categories ISL-2 inferrable? A corpus study
- **Satoru Ozaki, Rajesh Bhatt, Brian Dillon**. A LSTM language model learns Hindi-Urdu case-agreement interactions, and has a linear encoding of case

12-12:30: Methodology show-and-tell session

12:30-1:45: Lunch

1:45-2:45: Pragmatics Session 2

- **Sarah Ethridge, Joseph C. Y. Lau, Bronya R. Chernyak, Rob Voigt, Matthew Goldrick, Joseph Keshet, Molly Losh**. Self-Supervised Speech Representations in a Pre-train Speech Model Represent Key Rapid Automatized Naming Variability in Autism
- **Sanghee J. Kim, Forrest Davis**. Discourse Sensitivity in Attraction Effects: The Interplay Between Language Model Size and Training Data
- **Polina Tsvilodub, Robert D. Hawkins, Michael Franke**. Integrating Neural and Symbolic Components in a Model of Pragmatic Question-Answering

2:45-3:45: Poster session 2 (and coffee)

- **Thomas Graf, Kenneth Hanson**. Syntax with strings attached
- **Gustavo Cilleruelo Calderón, Emily Allaway, Barry Haddow, Alexandra Birch**. Generics are puzzling. Can language models find the missing piece?
- **Xiaomeng Zhu, Zhenghao Zhou, Simon Charlow, Robert Frank**. Do LLMs Understand Anaphoric Accessibility?
- **Dingyi Pan, Andrew Kehler**. Pragmatic Competence in LLMs: The Case of Eliciture
- **Vincent Czarnecki**. The Logic of Linearization: Interpretations of Trees via Strings
- **Andrew Liu, Gerald Penn**. Similarity, Transformation and the Newly Found Invariance of Influence Functions
- **Laurestine Bradford**. Aspectual classes as lexically-conditioned predictors of aspectual choice

3:45-4:45: Talk session 4 (Chair: Gaja Jarosz)

- **Yifan Hu**. BMRS-Net: Learning BMRS Predicates as Decision Trees

- **Connor Mayer**. Reconciling categorical and gradient models of phonotactics
- **Coleman Haley**. Unlocking finite-state morphological transducers: Derivational networks for Inuit-Yupik languages

4:45-5:45: Keynote 4: Daniel Fried

Improving NLP Systems via Speaker Listener Games

Across a diverse range of task-oriented settings, we've found that NLP models can interact more successfully with human partners by reasoning about language use as a strategic action. Building on the Rational Speech Acts framework, we formulate tasks as referential games played between a speaker and listener. We train speaker and listener models, implemented via grounded language models, and layer a reasoning procedure on top of them. This procedure uses the speaker/listener models to (1) predict how a person might interpret language from the system and (2) resolve ambiguity by reasoning counterfactually about what goal might have made a person say what they did. We show examples of this approach making interactions between people and systems more successful: in grounded instruction following, regular expression induction, and improving communicative efficiency by forming conceptual pacts.

5:45-6: Concluding remarks

6-6:30: Business meeting